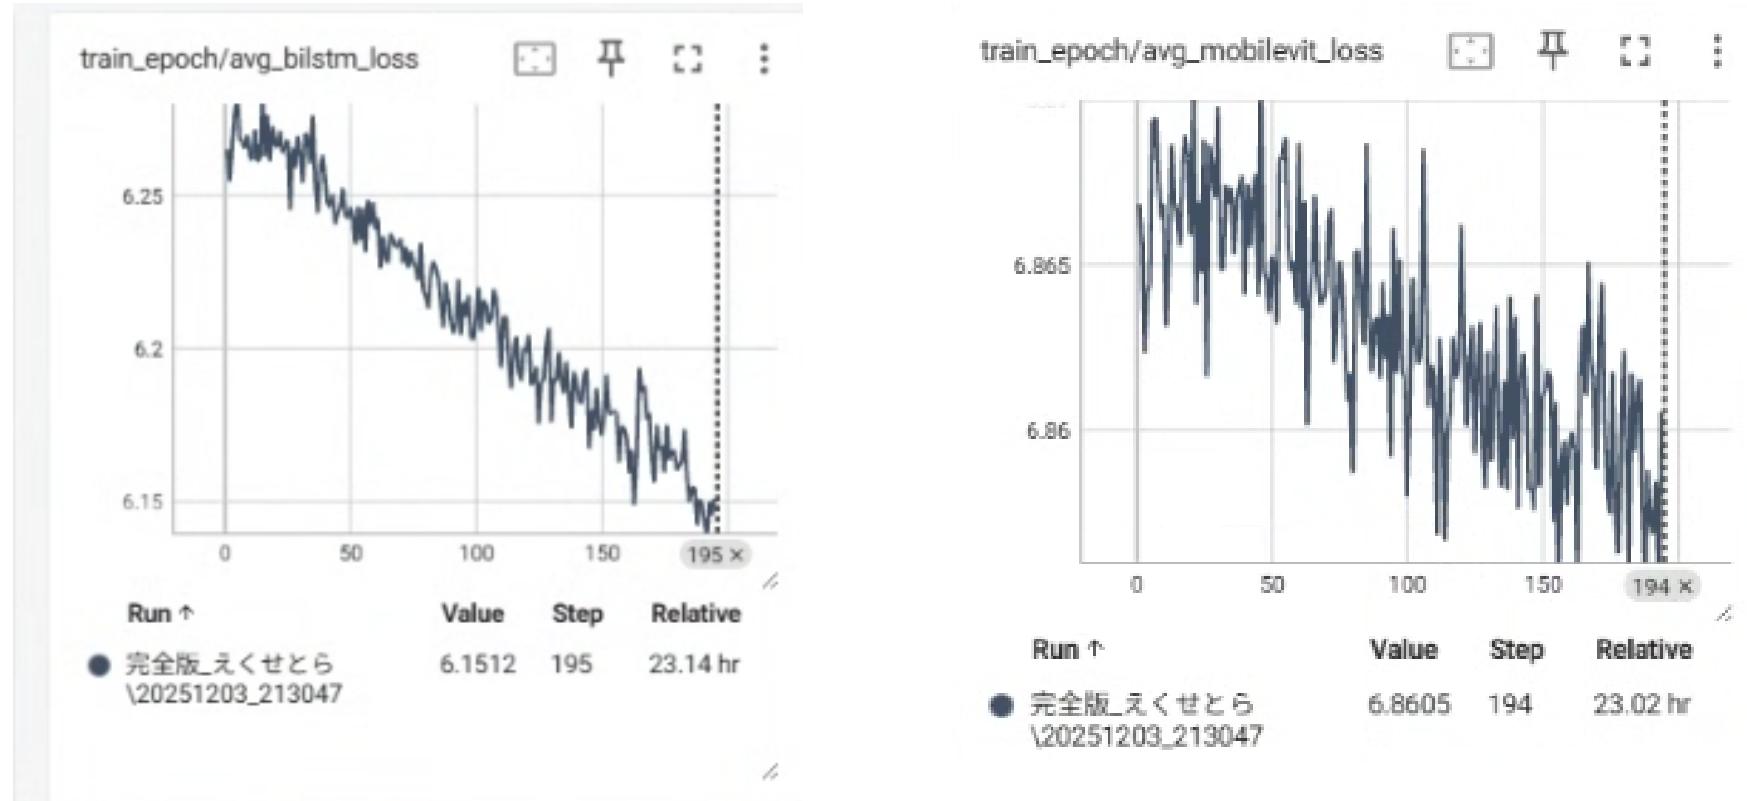
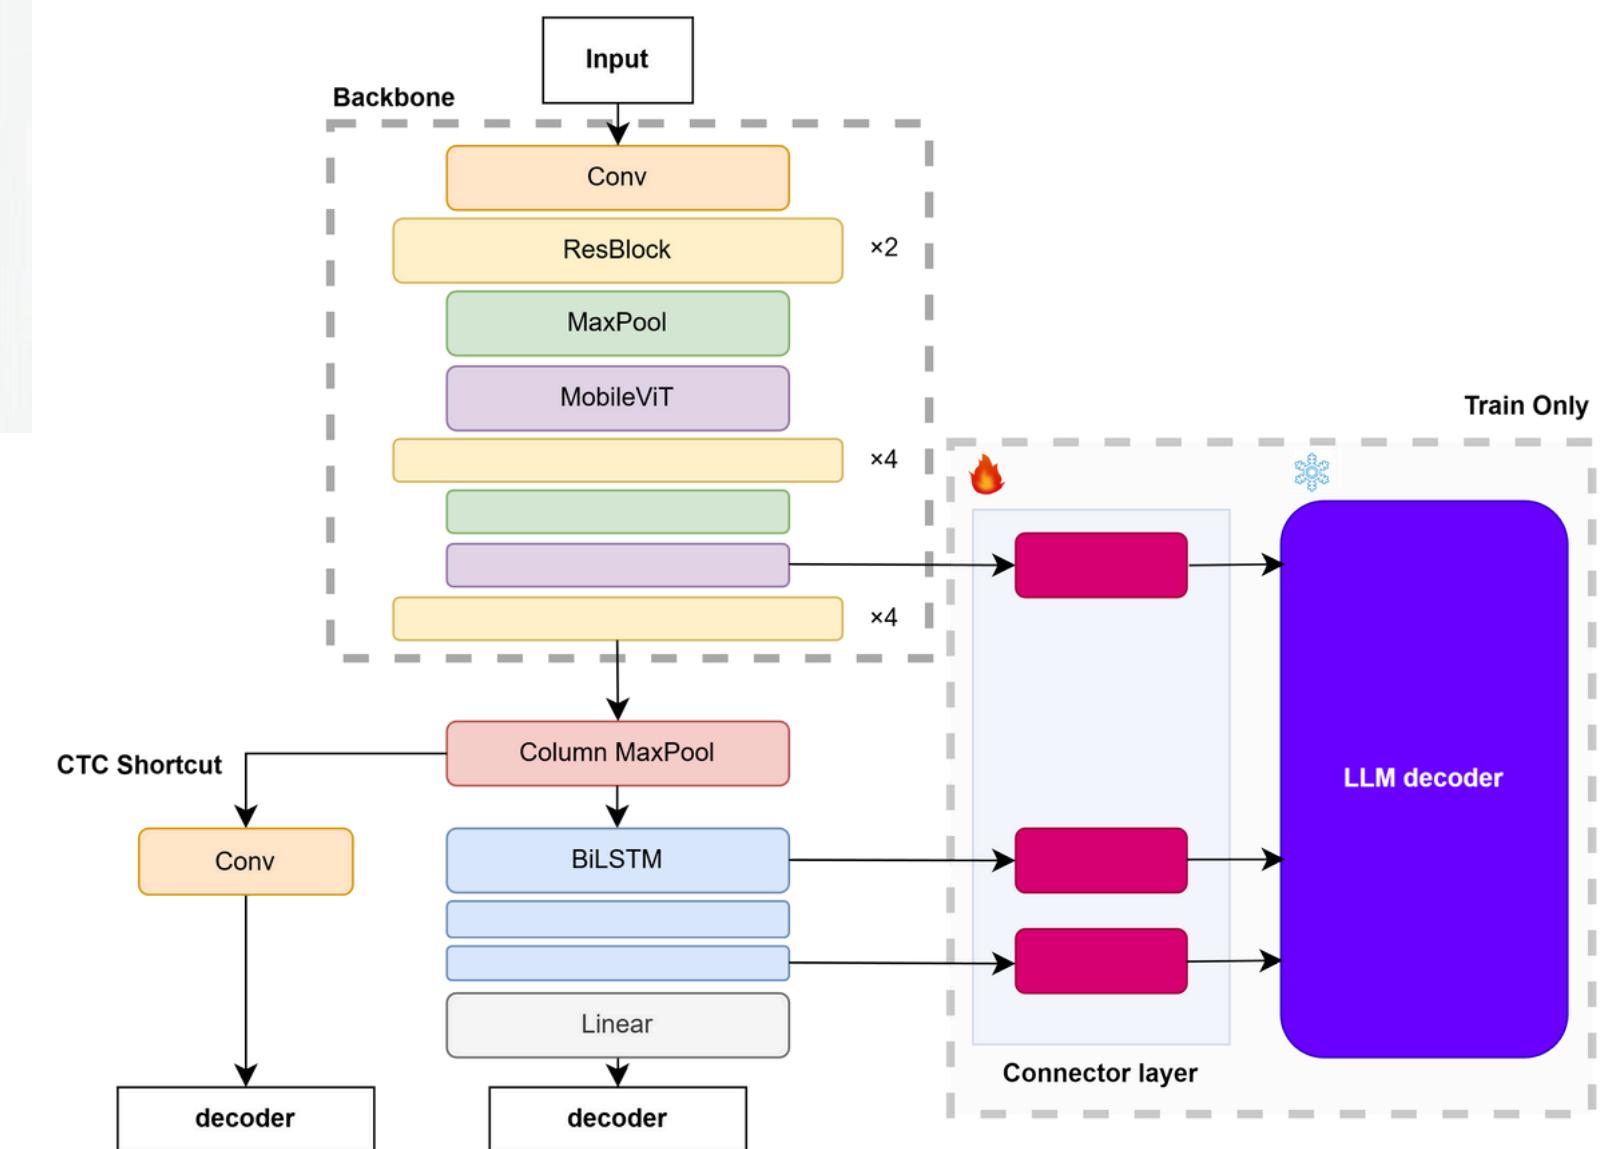


# 卒研にむけて②

# 提案手法前確認



lossの比較



前のモデル

# なにをしているのか

正しくファインチューニングできるのか（自分の工夫含まず）

→ファインチューニングできたら自分で書いた学習ループが正しいことが示せる

上記が確認できたら自分の工夫入れていこうと思う

# Trocrファインチューニング編

Model	Architecture	Training Data	External LM	CER
TrOCR <sub>SMALL</sub>	Transformer	Synthetic + IAM	No	4.22
TrOCR <sub>BASE</sub>	Transformer	Synthetic + IAM	No	3.42
TrOCR <sub>LARGE</sub>	Transformer	Synthetic + IAM	No	2.89

目標値

smallで精度確認→CER：70超

↓

≡ TrOCR (small-sized model, fine-tuned on IAM)

TrOCR model fine-tuned on the [IAM dataset](#). It was introduced in the paper [TrOCR: Transformer-based Optical Character Recognition with Pre-trained Models](#) by Li et al. and first released in [this repository](#).

学習済み重み

もうこれはすでに  
IAMでファインチューニング済  
なので同じデータとモデルを使っ  
てるおりかつ正しいモデル設定、  
データ前処理がなされていたら精  
度がでるはず

# Trocrファインチューニング編

## モデル設定見直し

```
model.config.pad_token_id = 1  
model.config.eos_token_id = 2  
model.config.decoder_start_token_id = 2
```

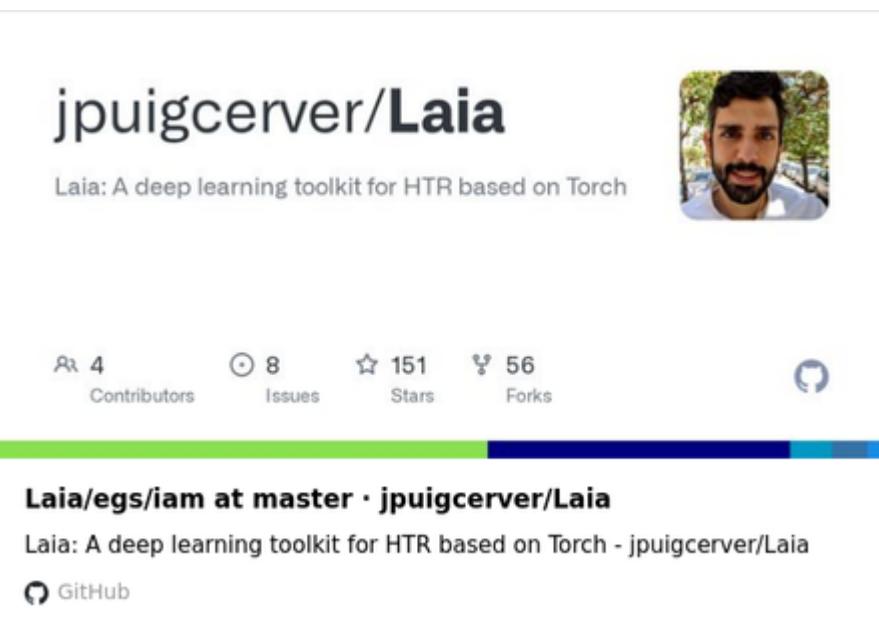
←CER:5.38

## 上記+データ見直し

- train,valのデータ分布
- リサイズ
- 論文記載のデータ処理を使用

←CER:4.33

(目標値にほぼ到達で提案手法に移れる！！)



# Trocrファインチューニング編

モデル設定見直し

```
model.config.pad_token_id = 1  
model.config.eos_token_id = 2  
model.config.decoder_start_token_id = 2
```

←CER:5.38

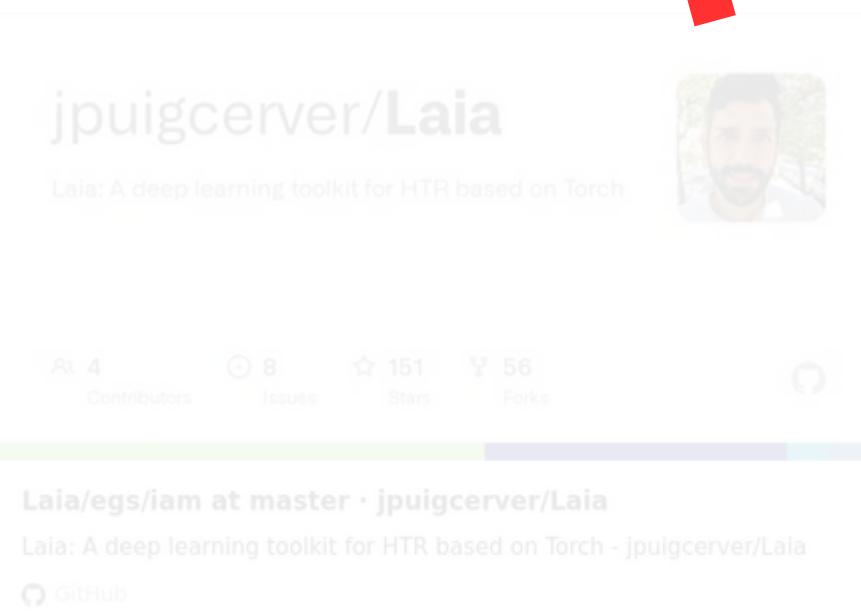
データ見直し

- train,valのデータ分離
- リサイズ
- 論文記載のデータ処理を使用

←CER:4.33

(目標値にほぼ到達で提案手法に移れる！！)

問題発生！！



# Trocrファインチューニング編

```
=====
TensorBoard logging to: ./logs/runs\20251208_205704_12-07_TOCR-small
Start TensorBoard with: python -m tensorboard --logdir="./logs/runs\20251208_205704_12-07_TOCR-small"
=====
Epoch 1: 0%| 0/1541 [00:00<?, ?it/s]
:\Users\user\AppData\Roaming\Python\Python310\site-packages\transformers\integrations\sdpa_attention.py:54: UserWarning: Torch was not compiled with flash attention. (Triggered internally at ..\aten\src\ATen\native\transformers\cuda\sdp_utils.cpp:263.)
    attn_output = torch.nn.functional.scaled_dot_product_attention(
Epoch 1: 100%| 1541/1541 [02:42<00:00, 9.49it/s, loss=2.9129]
Epoch 2: 100%| 1541/1541 [02:42<00:00, 9.51it/s, loss=1.6902]
Epoch 3: 100%| 1541/1541 [02:41<00:00, 9.52it/s, loss=1.3823]
Epoch 4: 100%| 1541/1541 [02:41<00:00, 9.53it/s, loss=1.1408]
Epoch 5: 100%| 1541/1541 [02:41<00:00, 9.54it/s, loss=1.0002]
Epoch 6: 100%| 1541/1541 [02:42<00:00, 9.51it/s, loss=0.8833]
Epoch 7: 100%| 1541/1541 [02:41<00:00, 9.53it/s, loss=0.7851]
Epoch 8: 100%| 1541/1541 [02:41<00:00, 9.54it/s, loss=0.7385]
Epoch 9: 100%| 1541/1541 [02:42<00:00, 9.51it/s, loss=0.6279]
Epoch 10: 100%| 1541/1541 [02:41<00:00, 9.52it/s, loss=0.5730]
Test epoch 10: 100%| 365/365 [01:57<00:00, 3.10it/s]
Test CER: 0.3085
```

ログ

# Trocrファインチューニング編

## 課題

学習を回すと精度が悪化する

- 現状：epoch:10、CER: 30.85
- 違和感：ロスは減少している、まず、valとtestで差が10倍

## 試したこと

- 1部の層のみ学習可能にする (encoderの最終層とdecoder 2 層) → CER : 5.90
- ロスを自作 (内部で何してるかわからん！) → CER : 8.06
- 上二つを組み合わせた：CER : 5.16

```
Test epoch 10: 100% | 0/1541 [00:00, ?it/s]C
Test CER: 0.3085
TensorBoard logging to: ./logs/runs\20251208_212930_12-07_TOCR-small
Start TensorBoard with: python -m tensorboard --logdir=.\logs\runs\20251208_212930_12-07_TOCR-small
=====
Epoch 1: 100%
:Users\user\AppData\Roaming\Python\Python310\site-packages\transformers\integrations\sdp_attention.py:54: UserWarning: lTorch was not compiled with flash attention. (Triggered internally at ...aten\src\Aten\native\transformers\cuda\sdp_utils.cpp:263 )
    attn_output = torch.nn.functional.scaled_dot_product_attention(
=====
| 0/1541 [00:00<?, ?it/s]
1541/1541 [01:14&lt;00:00, 20.64it/s, loss=0.6422]
1541/1541 [01:15&lt;00:00, 20.45it/s, loss=0.1520]
1541/1541 [01:15&lt;00:00, 20.53it/s, loss=0.1240]
1541/1541 [01:14&lt;00:00, 20.66it/s, loss=0.1063]
1541/1541 [01:14&lt;00:00, 20.59it/s, loss=0.0939]
1541/1541 [01:15&lt;00:00, 20.54it/s, loss=0.0851]
1541/1541 [01:14&lt;00:00, 20.57it/s, loss=0.0798]
1541/1541 [01:15&lt;00:00, 20.55it/s, loss=0.0692]
1541/1541 [01:15&lt;00:00, 20.67it/s, loss=0.0657]
1541/1541 [01:14&lt;00:00, 20.78it/s, loss=0.0621]
| 365/365 [01:46&lt;00:00, 3.43it/s]

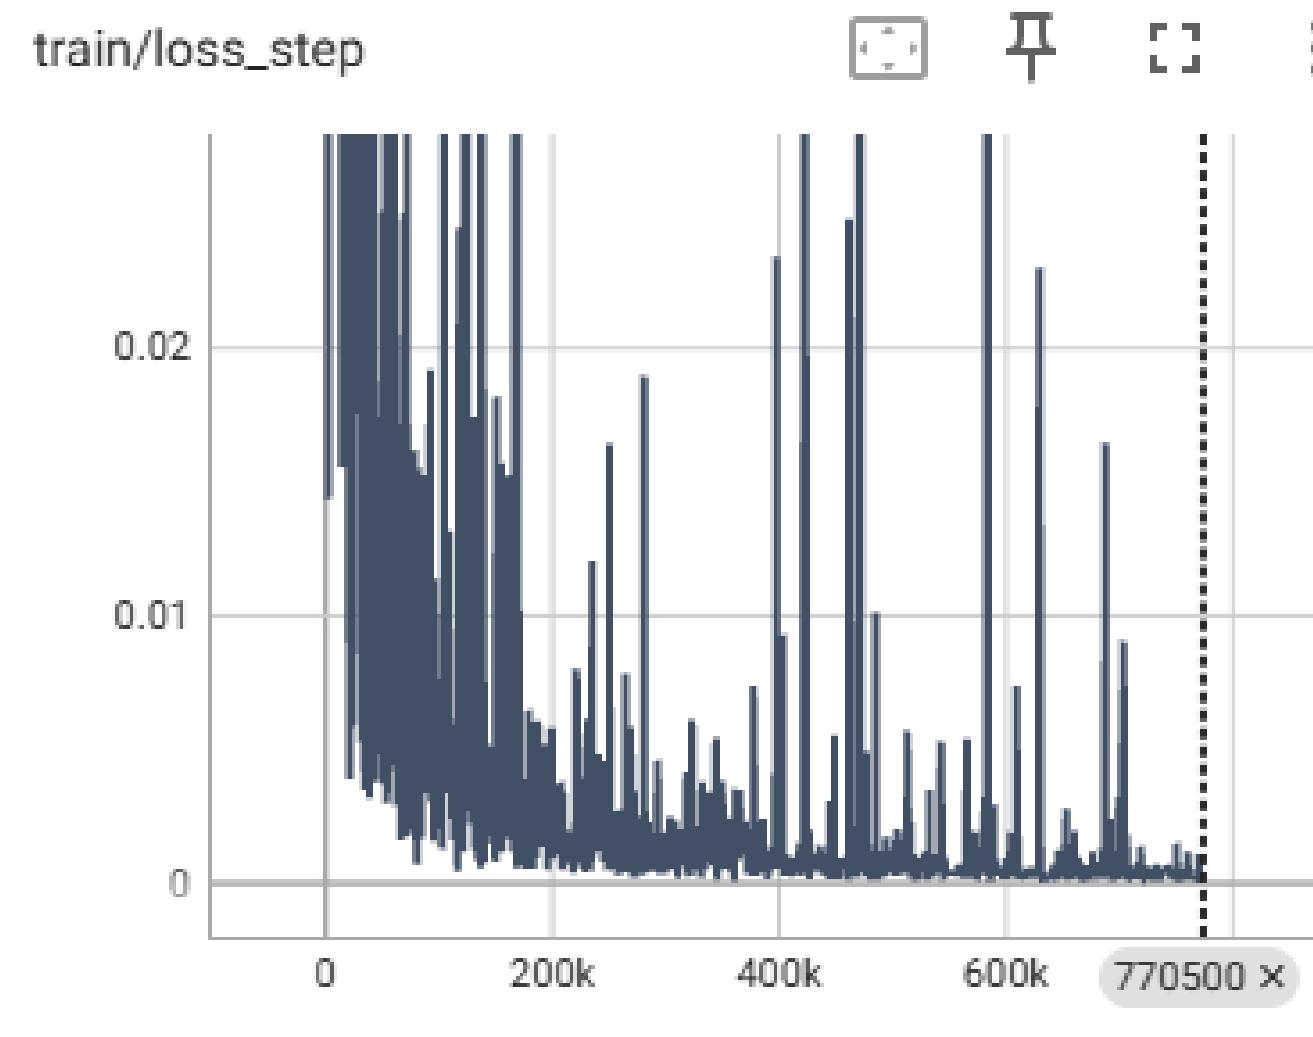
Test epoch 10: 100% | 0/1541 [00:00<?, ?it/s]C
Test CER: 0.0517</pre
```

学習なしの評価→**CER:4.33**でいずれも精度悪化

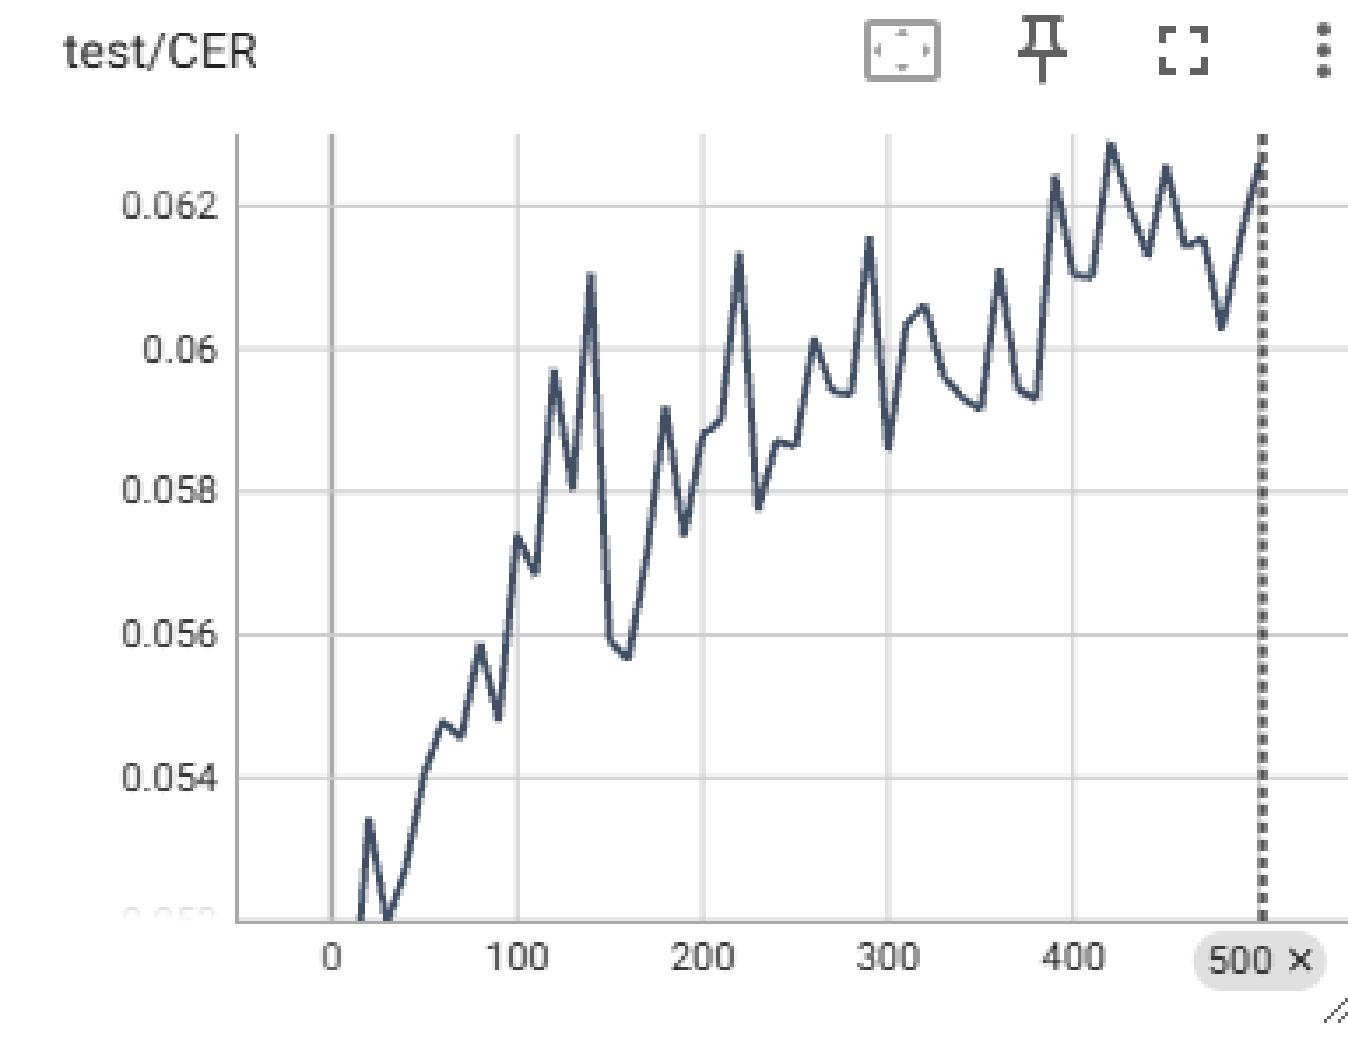
lossは下がってるしもう少し学習ステップを増やしてみる→

# Trocrファインチューニング編

## ログ



Run ↑	Value	Step	Relative
20251209_151417_12-07_TOCR-small	0	770,500	14.92 hr



Run ↑	Value	Step	Relative
20251209_151417_12-07_TOCR-small	0.0626	500	14.66 hr

ステップごとの値とってしまった見にくいですよね...lol

# モデル構造

```
(encoder): ViTModel(
    (embeddings): ViTEmbeddings(
        (patch_embeddings): ViTPatchEmbeddings(
            (projection): Conv2d(3, 768, kernel_size=(16, 16), stride=(16, 16))
        )
        (dropout): Dropout(p=0.0, inplace=False)
    )
    (encoder): ViTEncoder(
        (layer): ModuleList(
            (0-11): 12 x ViTLayer(
                (attention): ViTAttention(
                    (attention): ViTSelfAttention(
                        (query): Linear(in_features=768, out_features=768, bias=False)
                        (key): Linear(in_features=768, out_features=768, bias=False)
                        (value): Linear(in_features=768, out_features=768, bias=False)
                    )
                    (output): ViTSelfOutput(
                        (dense): Linear(in_features=768, out_features=768, bias=True)
                        (dropout): Dropout(p=0.0, inplace=False)
                    )
                )
                (intermediate): ViTIntermediate(
                    (dense): Linear(in_features=768, out_features=3072, bias=True)
                    (intermediate_act_fn): GELUActivation()
                )
                (output): ViTOOutput(
                    (dense): Linear(in_features=3072, out_features=768, bias=True)
                    (dropout): Dropout(p=0.0, inplace=False)
                )
                (layernorm_before): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
                (layernorm_after): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            )
        )
        (layernorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (pooler): ViTPooler(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (activation): Tanh()
        )
    )
)
```

```
(decoder): TrOCRForCausalLM(
    (model): TrOCRDecoderWrapper(
        (decoder): TrOCRDecoder(
            (embed_tokens): TrOCRScaledWordEmbedding(50265, 1024, padding_idx=1)
            (embed_positions): TrOCRLearnedPositionalEmbedding(514, 1024)
            (layernorm_embedding): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
            (layers): ModuleList(
                (0-11): 12 x TrOCRDecoderLayer(
                    (self_attn): TrOCRAttention(
                        (k_proj): Linear(in_features=1024, out_features=1024, bias=True)
                        (v_proj): Linear(in_features=1024, out_features=1024, bias=True)
                        (q_proj): Linear(in_features=1024, out_features=1024, bias=True)
                        (out_proj): Linear(in_features=1024, out_features=1024, bias=True)
                    )
                    (activation_fn): GELUActivation()
                    (self_attn_layer_norm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
                    (encoder_attn): TrOCRAttention(
                        (k_proj): Linear(in_features=768, out_features=1024, bias=True)
                        (v_proj): Linear(in_features=768, out_features=1024, bias=True)
                        (q_proj): Linear(in_features=1024, out_features=1024, bias=True)
                        (out_proj): Linear(in_features=1024, out_features=1024, bias=True)
                    )
                    (encoder_attn_layer_norm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
                    (fc1): Linear(in_features=1024, out_features=4096, bias=True)
                    (fc2): Linear(in_features=4096, out_features=1024, bias=True)
                    (final_layer_norm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
                )
            )
            (output_projection): Linear(in_features=1024, out_features=50265, bias=False)
        )
    )
)' loaded.
```